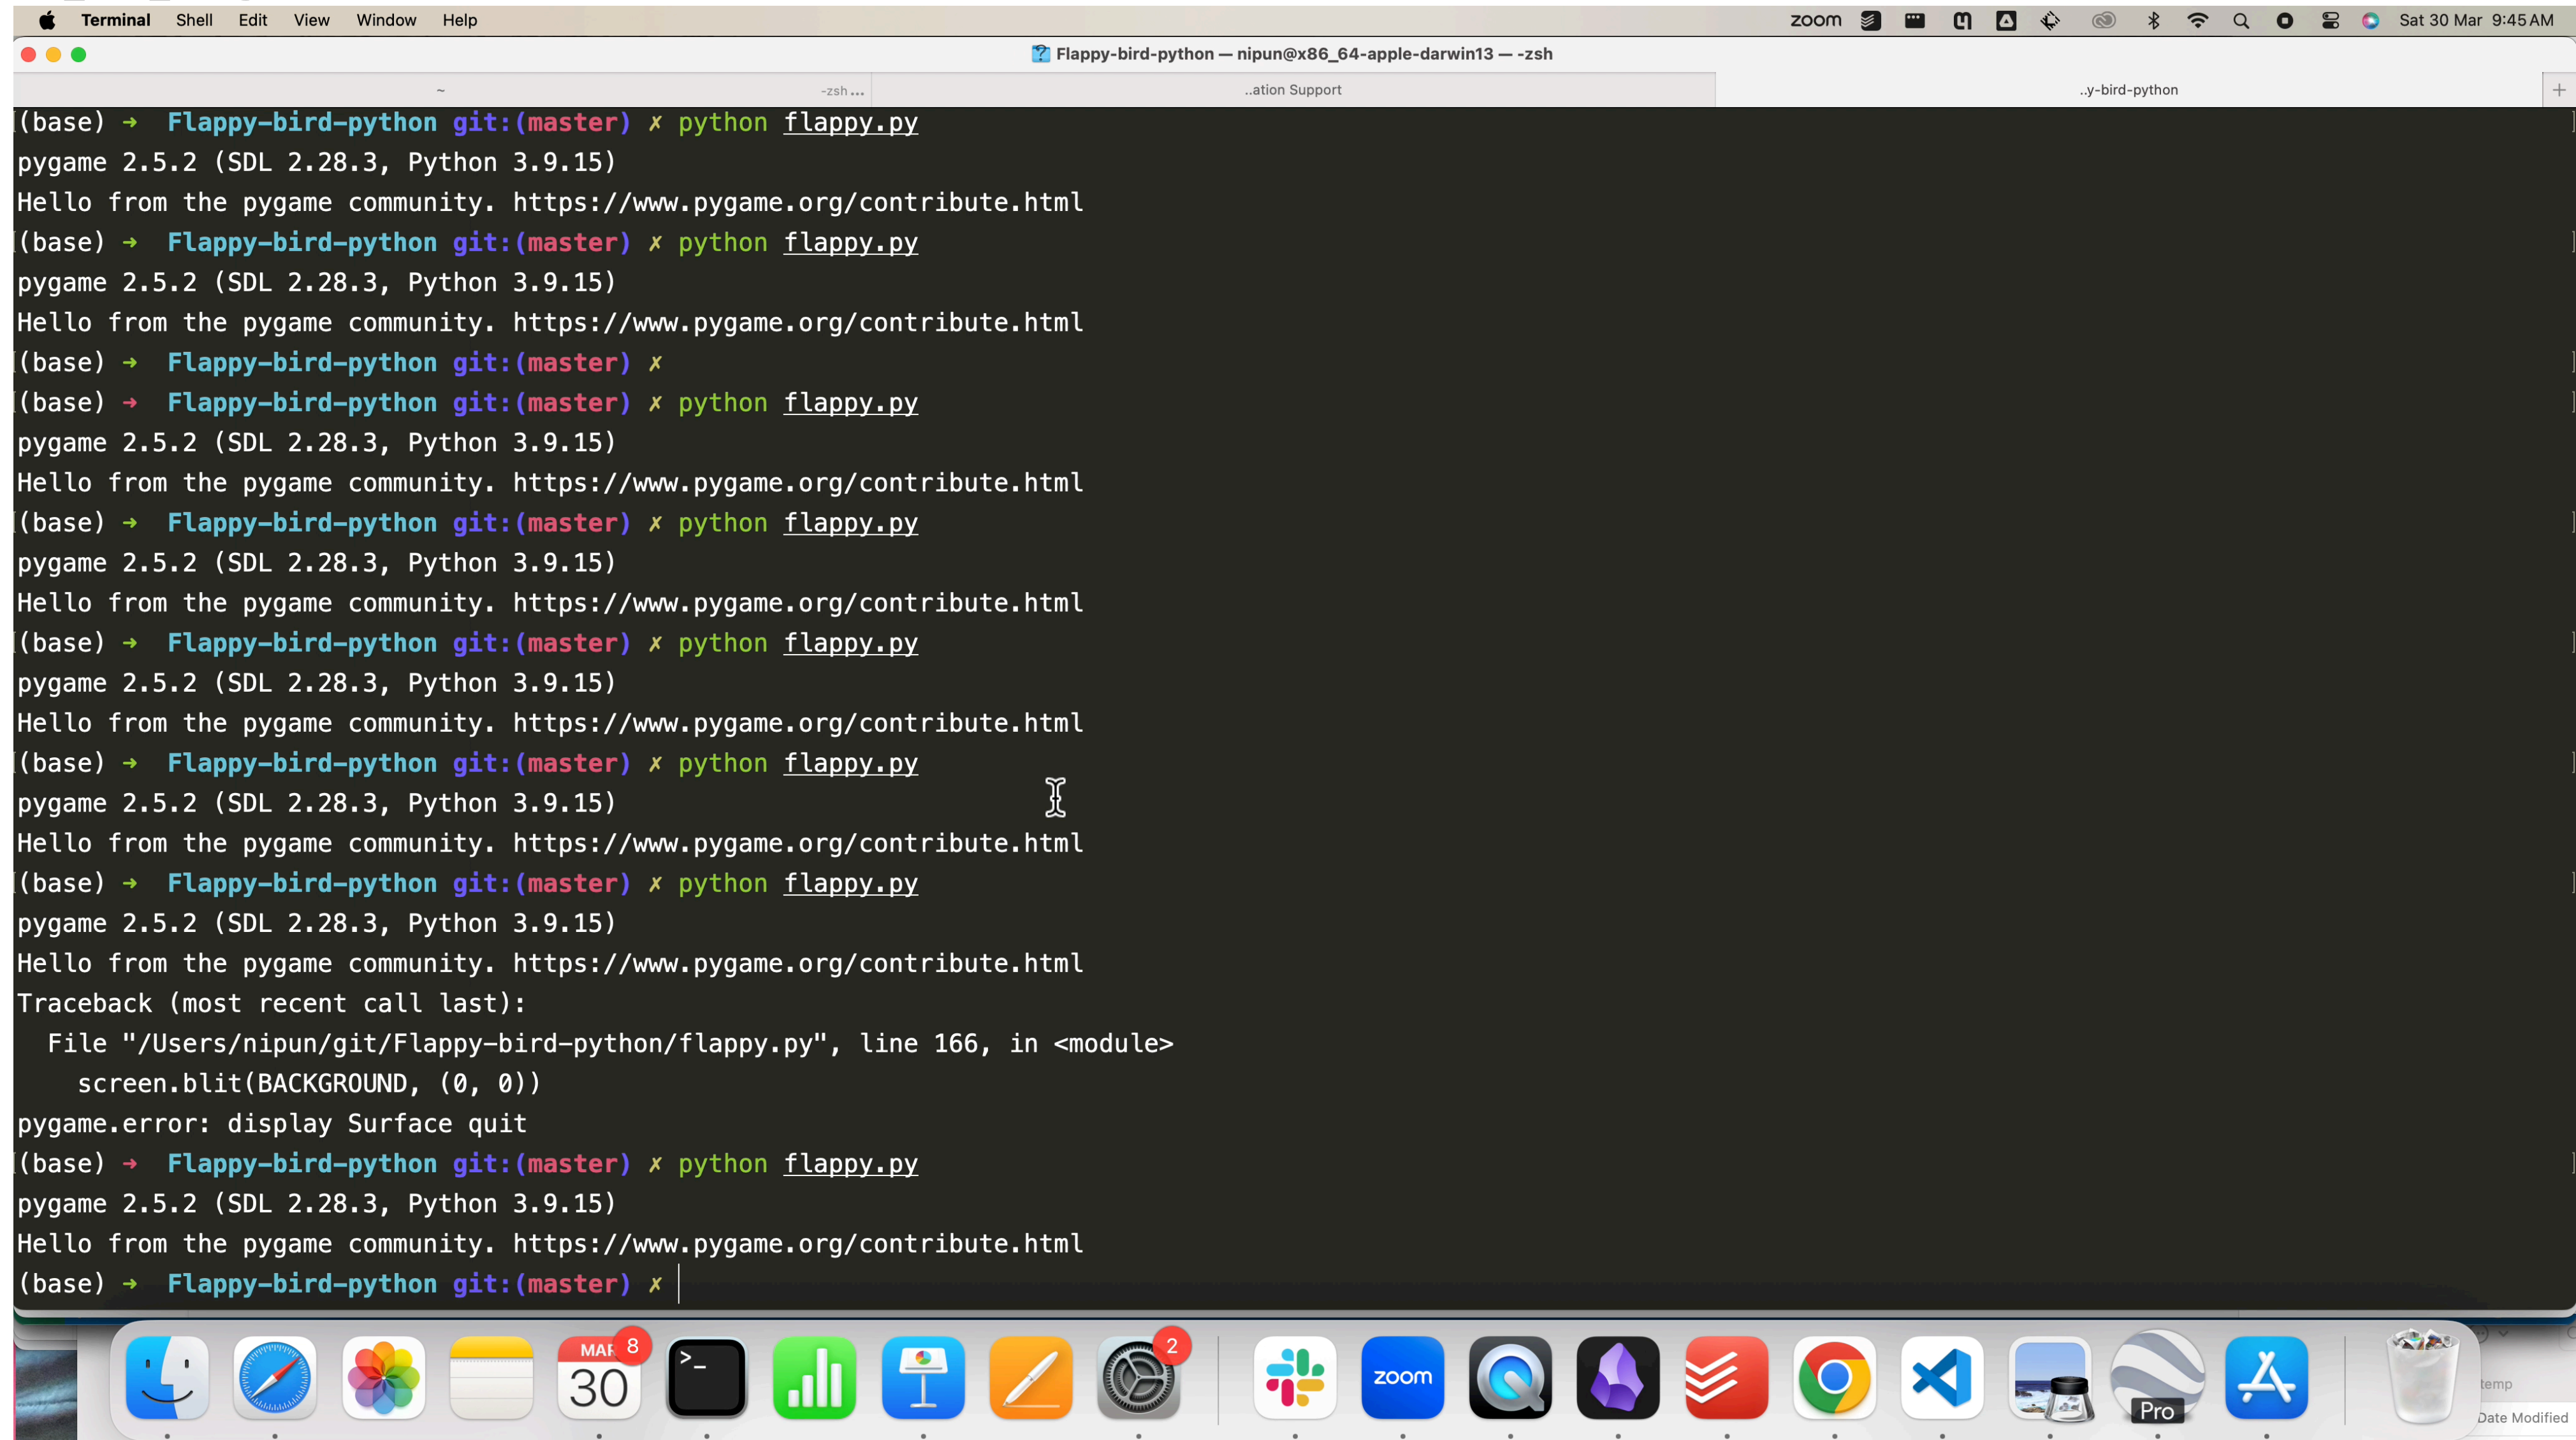


# Reinforcement Learning

Nipun Batra, 1 April 2024

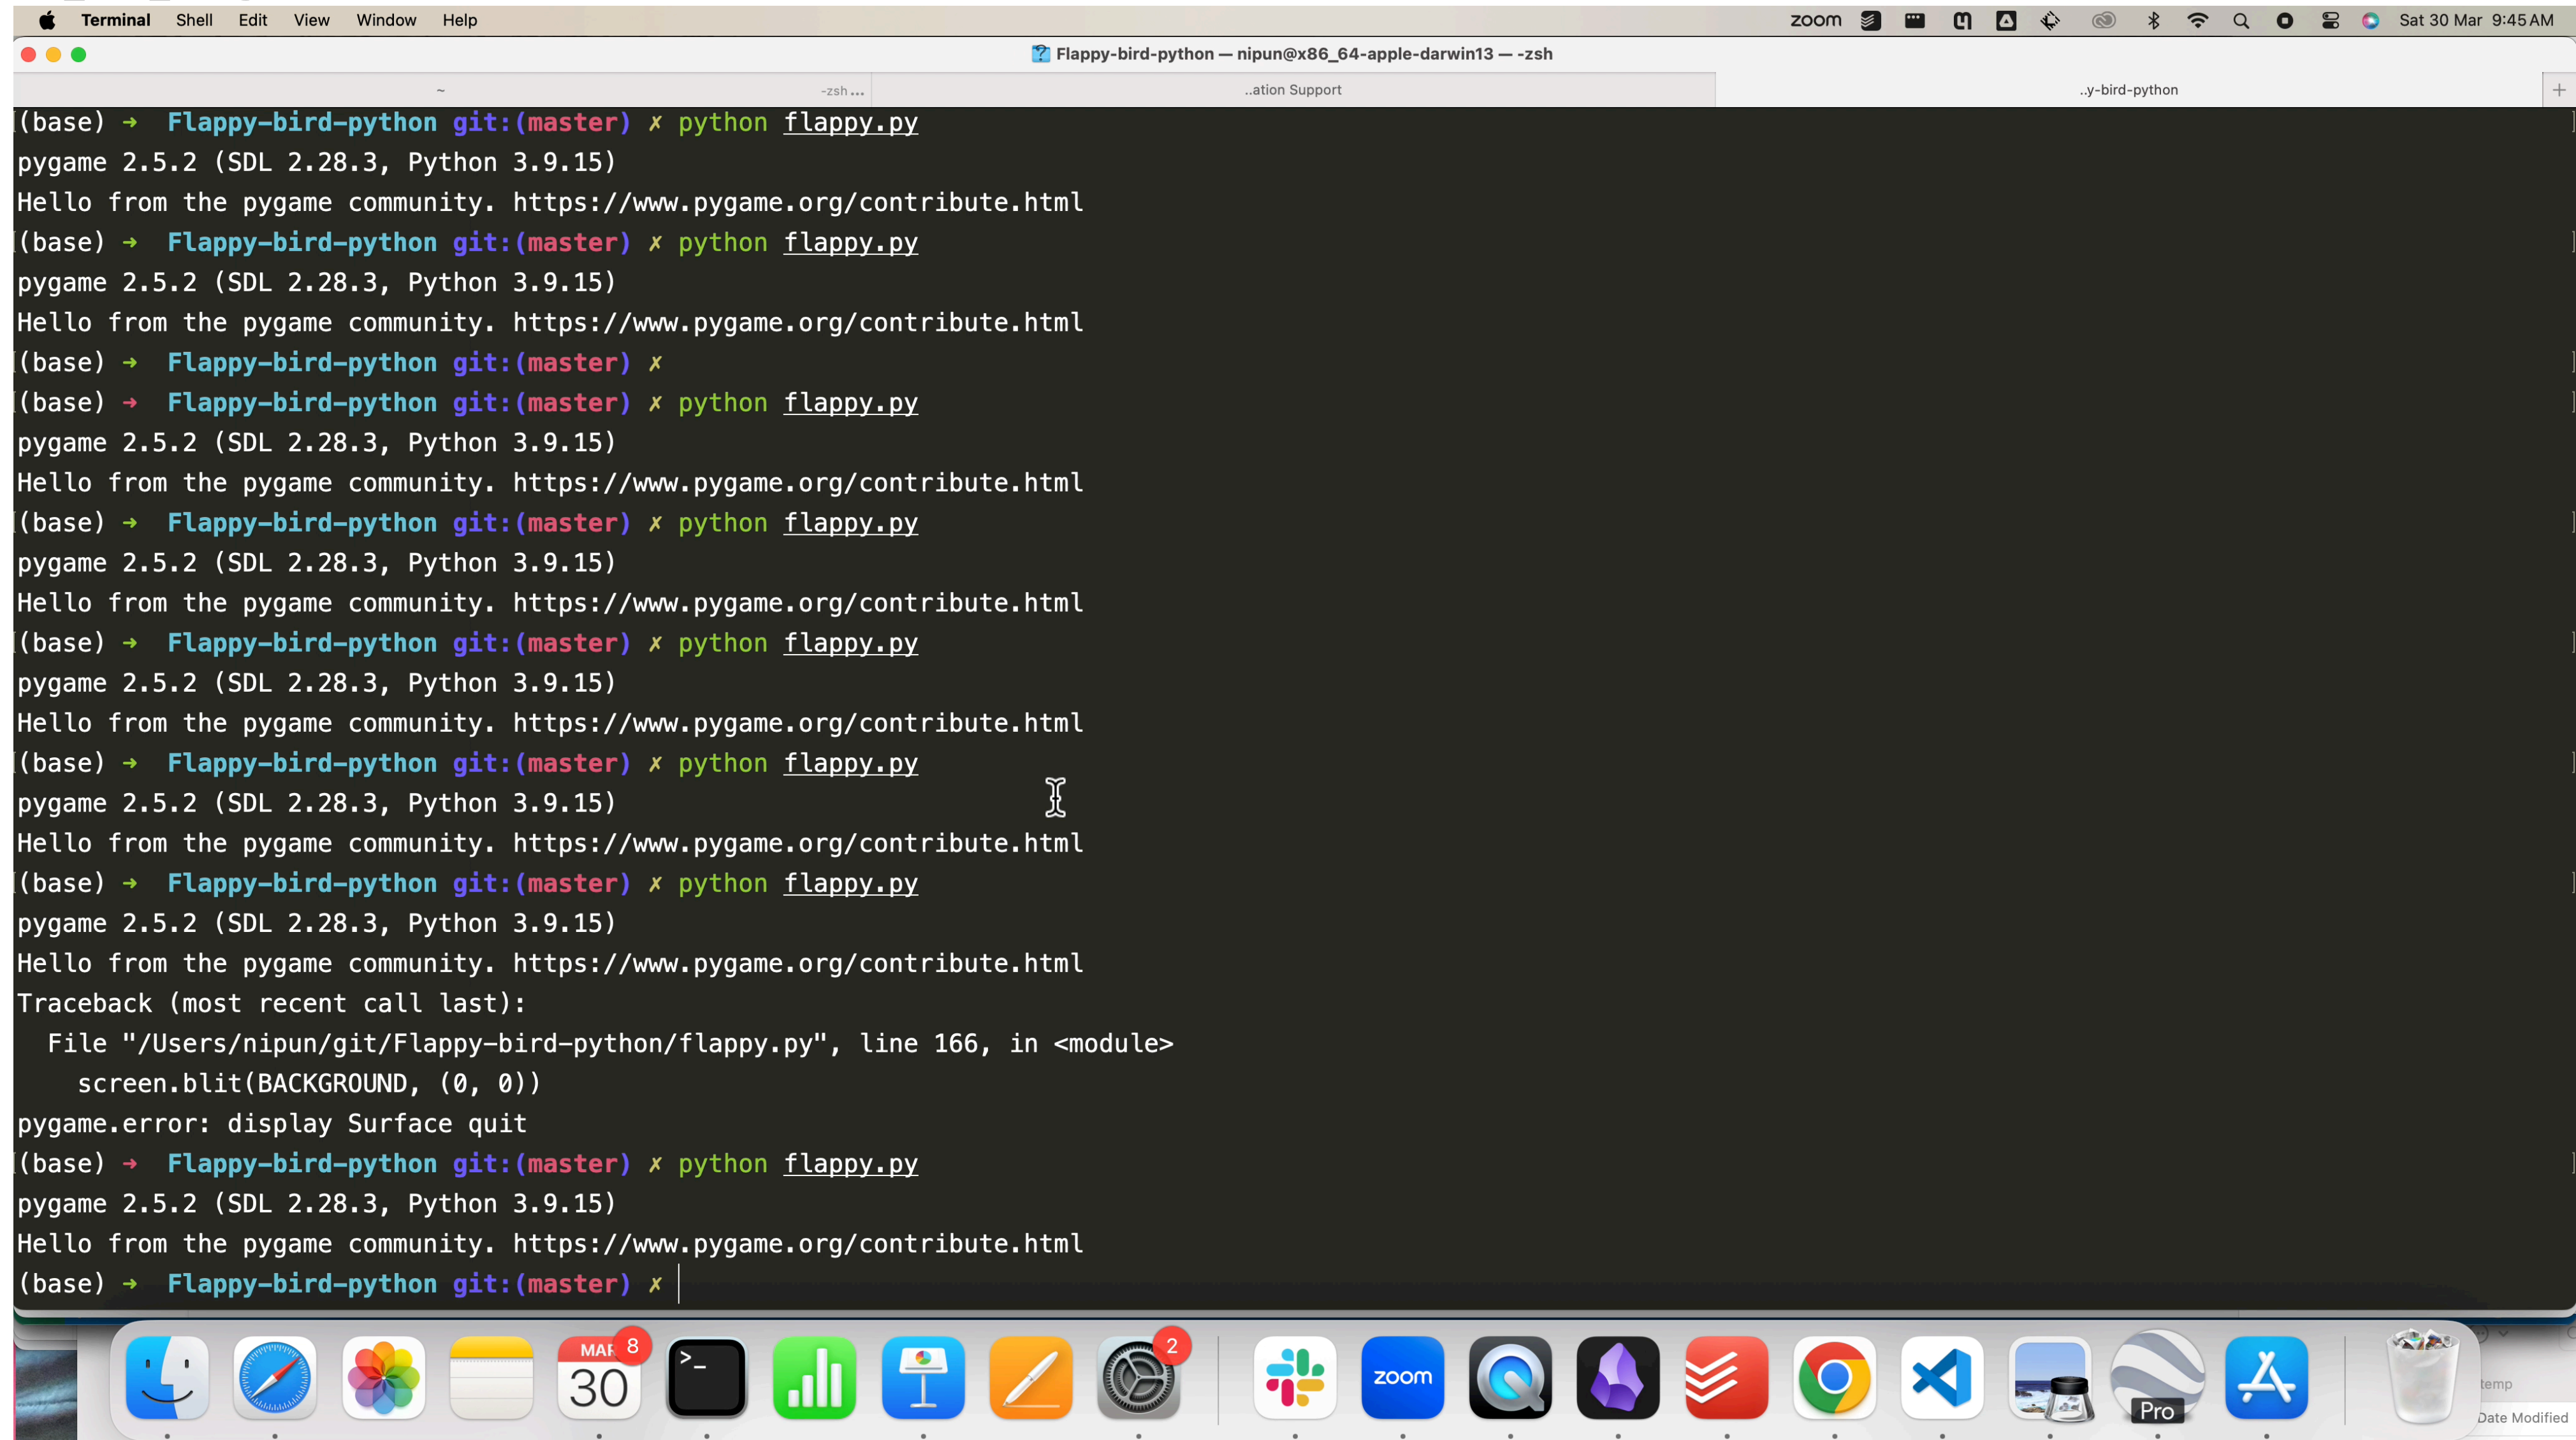
# Flappy Bird



```
(base) → Flappy-bird-python git:(master) x python flappy.py
pygame 2.5.2 (SDL 2.28.3, Python 3.9.15)
Hello from the pygame community. https://www.pygame.org/contribute.html
(base) → Flappy-bird-python git:(master) x python flappy.py
pygame 2.5.2 (SDL 2.28.3, Python 3.9.15)
Hello from the pygame community. https://www.pygame.org/contribute.html
(base) → Flappy-bird-python git:(master) x
(base) → Flappy-bird-python git:(master) x python flappy.py
pygame 2.5.2 (SDL 2.28.3, Python 3.9.15)
Hello from the pygame community. https://www.pygame.org/contribute.html
(base) → Flappy-bird-python git:(master) x python flappy.py
pygame 2.5.2 (SDL 2.28.3, Python 3.9.15)
Hello from the pygame community. https://www.pygame.org/contribute.html
(base) → Flappy-bird-python git:(master) x python flappy.py
pygame 2.5.2 (SDL 2.28.3, Python 3.9.15)
Hello from the pygame community. https://www.pygame.org/contribute.html
(base) → Flappy-bird-python git:(master) x python flappy.py
pygame 2.5.2 (SDL 2.28.3, Python 3.9.15)
Hello from the pygame community. https://www.pygame.org/contribute.html
Traceback (most recent call last):
  File "/Users/nipun/git/Flappy-bird-python/flappy.py", line 166, in <module>
    screen.blit(BACKGROUND, (0, 0))
pygame.error: display Surface quit
(base) → Flappy-bird-python git:(master) x python flappy.py
pygame 2.5.2 (SDL 2.28.3, Python 3.9.15)
Hello from the pygame community. https://www.pygame.org/contribute.html
(base) → Flappy-bird-python git:(master) x
```



# Flappy Bird




```
(base) → Flappy-bird-python git:(master) x python flappy.py
pygame 2.5.2 (SDL 2.28.3, Python 3.9.15)
Hello from the pygame community. https://www.pygame.org/contribute.html
(base) → Flappy-bird-python git:(master) x python flappy.py
pygame 2.5.2 (SDL 2.28.3, Python 3.9.15)
Hello from the pygame community. https://www.pygame.org/contribute.html
(base) → Flappy-bird-python git:(master) x
(base) → Flappy-bird-python git:(master) x python flappy.py
pygame 2.5.2 (SDL 2.28.3, Python 3.9.15)
Hello from the pygame community. https://www.pygame.org/contribute.html
(base) → Flappy-bird-python git:(master) x python flappy.py
pygame 2.5.2 (SDL 2.28.3, Python 3.9.15)
Hello from the pygame community. https://www.pygame.org/contribute.html
(base) → Flappy-bird-python git:(master) x python flappy.py
pygame 2.5.2 (SDL 2.28.3, Python 3.9.15)
Hello from the pygame community. https://www.pygame.org/contribute.html
(base) → Flappy-bird-python git:(master) x python flappy.py
pygame 2.5.2 (SDL 2.28.3, Python 3.9.15)
Hello from the pygame community. https://www.pygame.org/contribute.html
Traceback (most recent call last):
  File "/Users/nipun/git/Flappy-bird-python/flappy.py", line 166, in <module>
    screen.blit(BACKGROUND, (0, 0))
pygame.error: display Surface quit
(base) → Flappy-bird-python git:(master) x python flappy.py
pygame 2.5.2 (SDL 2.28.3, Python 3.9.15)
Hello from the pygame community. https://www.pygame.org/contribute.html
(base) → Flappy-bird-python git:(master) x
```

# Flappy Bird

- Game demo (Code modified from: <https://github.com/LeonMarqs/Flappy-bird-python>)

# Flappy Bird


What is the goal/objective?



# Flappy Bird


What is the goal/objective?

- Maximise score



# Flappy Bird


What are the actions we can take?



# Flappy Bird

What are the actions we can take?


- Two actions
  - Tap (Space)
  - No tap





# Flappy Bird

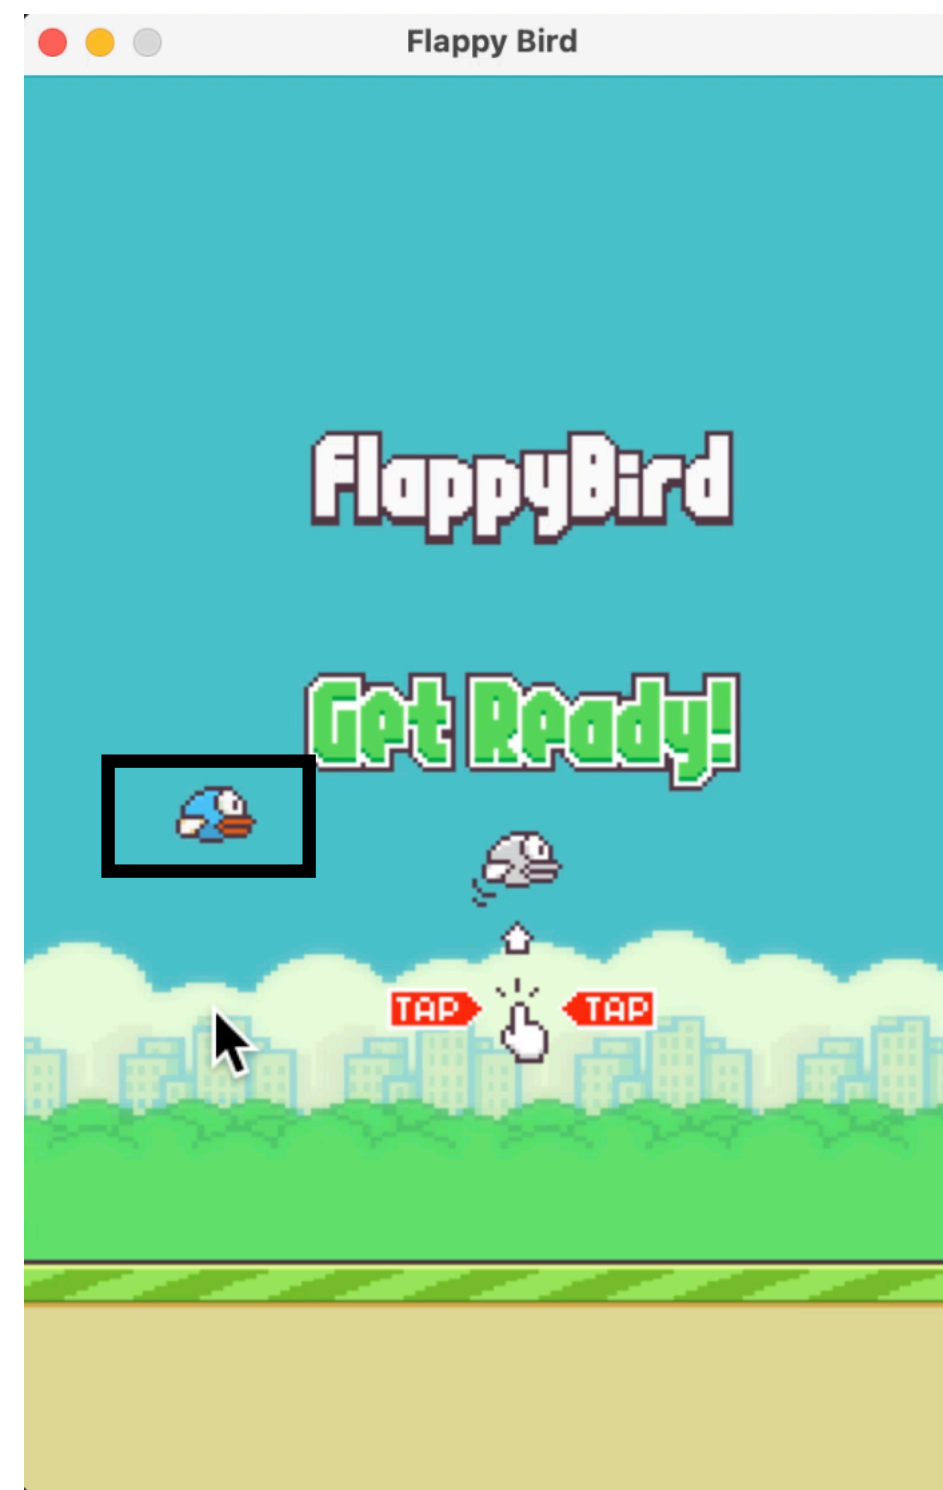
## Who is playing?



# Flappy Bird

## Who is playing?

- Agent
  - You
  - Or some algorithm



# Flappy Bird

## Where are we playing?

```
while begin:

    clock.tick(15)

    for event in pygame.event.get():
        if event.type == QUIT:
            pygame.quit()
        if event.type == KEYDOWN:
            if event.key == K_SPACE or event.key == K_UP:
                bird.bump()
                pygame.mixer.music.load(wing)
                pygame.mixer.music.play()
                begin = False


    screen.blit(BACKGROUND, (0, 0))
    screen.blit(BEGIN_IMAGE, (120, 150))

    if is_off_screen(ground_group.sprites()[0]):
        ground_group.remove(ground_group.sprites()[0])

        new_ground = Ground(GROUND_WIDHT - 20)
        ground_group.add(new_ground)

    bird.begin()
    ground_group.update()

    bird_group.draw(screen)
    ground_group.draw(screen)
```



# Flappy Bird

## Where are we playing?

- Environment
- Code
  - generating the graphics
  - Physics rules
    - What happens when you tap
    - What happens when you hit pipe

```
while begin:

    clock.tick(15)

    for event in pygame.event.get():
        if event.type == QUIT:
            pygame.quit()
        if event.type == KEYDOWN:
            if event.key == K_SPACE or event.key == K_UP:
                bird.bump()
                pygame.mixer.music.load(wing)
                pygame.mixer.music.play()
                begin = False


    screen.blit(BACKGROUND, (0, 0))
    screen.blit(BEGIN_IMAGE, (120, 150))

    if is_off_screen(ground_group.sprites()[0]):
        ground_group.remove(ground_group.sprites()[0])

        new_ground = Ground(GROUND_WIDHT - 20)
        ground_group.add(new_ground)

    bird.begin()
    ground_group.update()

    bird_group.draw(screen)
    ground_group.draw(screen)
```



# Flappy Bird


**What does the environment provide to an agent?**



# Flappy Bird

What does the environment provide to an agent?

Reward




# Flappy Bird

What does the environment provide to an agent?

**Observations**


Pixel level information



# Flappy Bird

How does an agent decide what action to take?

Should the agent  
tap or not?




# Flappy Bird

How does an agent decide what action to take?

**State**

Process  
observation into a  
“state”




# Flappy Bird

How does an agent decide what action to take?

**State**


Process  
observation into a  
“state”





# Flappy Bird

## Is time important?




# Flappy Bird

## Is time important?

**Yes**


Agent's current  
state depends on  
previous state and  
action




# Block Diagram



# Block Diagram



# Block Diagram






# OpenAI Gym Environment

- Mountain Car
  - Actions?
  - State?

# Block Diagram



# Goal: Maximise total (discounted) reward

- Total Reward (Return)  $R_t = \sum_{i=t}^{\infty} r_i = r_t + r_{t+1} \dots + r_{t+n} + \dots$
- Total Reward (Discounted Return)  
$$R_t = \sum_{i=t}^{\infty} \gamma^i r_i = \gamma^t r_t + \gamma^{t+1} r_{t+1} \dots + \gamma^{t+n} r_{t+n} + \dots$$
- $\gamma$  : discount factor;  $0 < \gamma < 1$

# Q function

- What we want?
  - Given a state choose an “action” that maximises total discounted reward
- Total Reward (Discounted Return)
$$R_t = \sum_{i=t}^{\infty} \gamma^i r_i = \gamma^t r_t + \gamma^{t+1} r_{t+1} \dots + \gamma^{t+n} r_{t+n} + \dots$$
- $Q(s_t, a_t) = \mathbb{E}[R_t \mid s_t, a_t]$
- Q-function captures the expected total future reward an agent can achieve by taking an action.

# Q table

State	Action 1	Action 2	Action 3
S1	10	20	15
S2	20	30	5
...			
SN	-5	10	20



# Q table

State	Action 1	Action 2	Action 3
S1	10	20	15
S2	20	30	5
...			
SN	-5	10	20

What action will you choose if you are in state S2?

# Q table

State	Action 1	Action 2	Action 3
S1	10	20	15
S2	20	30	5
...			
SN	-5	10	20

What action will you choose if you are in state S2?

Action 2 (as it gives us highest return)

# Q table

State <Position, Velocity>	Action 1	Action 2	Action 3
<-5, -2>	?	?	?
...	?	?	?
...			
...	?	?	?

How do we define states for problems like Mountain car where these numbers are not discrete?

# Q table

State <Position, Velocity>	Action 1	Action 2	Action 3
<[-5, -4], [-2, -1]>	?	?	?
<[-5, -4], [-1, 0]>	?	?	?
...			
...	?	?	?

How do we define states for problems like Mountain car where these numbers are not discrete?

Discretisation (notebook)

# Q function (revision)

- What we want?
  - Given a state choose an “action” that maximises total discounted reward
- Total Reward (Discounted Return)
$$R_t = \sum_{i=t}^{\infty} \gamma^i r_i = \gamma^t r_t + \gamma^{t+1} r_{t+1} \dots + \gamma^{t+n} r_{t+n} + \dots$$
- $Q(s_t, a_t) = \mathbb{E}[R_t \mid s_t, a_t]$
- Q-function captures the expected total future reward an agent can achieve by taking an action.

# Bellman Equation

The Bellman equation for Q-values is given by:

$$Q(s, a) = R(s, a) + \gamma \cdot \max_{a'} Q(s', a')$$

where:

- $Q(s, a)$  is the Q-value of taking action
- $R(s, a)$  is the immediate reward of taking action  $a$  in state  $s$ .
- $\gamma$  is the discount factor that determines the importance of future rewards.
- $s'$  is the next state after taking action  $a$ .
- $\max_{a'} Q(s', a')$  is the maximum Q-value over all possible actions in state  $s'$ .

# Q-learning Update Bellman Equation

$$Q(s, a) = R(s, a) + \gamma \cdot \max_{a'} Q(s', a')$$

Q-learning update rule is derived by using the Bellman equation in an iterative manner:

$$Q(s, a) \leftarrow Q(s, a) + \alpha \cdot \left( R(s, a) + \gamma \cdot \max_{a'} Q(s', a') - Q(s, a) \right)$$

- $\alpha$  is the learning rate that controls the extent to which new information overrides old information.

- $R(s, a) + \gamma \cdot \max_{a'} Q(s', a') - Q(s, a)$  is the temporal difference (TD) error, representing the discrepancy between the expected Q-value and the observed reward.